



# Predictive Modeling

Pradeep Kumar Mishra

PGP-DSBA Online

Jun\_B\_21

Date: 28:Nov:2021

## Content View

<b>Problem Statement - 1 Linear Regression</b>	<b>3</b>
1.1. Read the data and do exploratory data analysis. Describe the data briefly. (Check the null values, Data types, shape, EDA, duplicate values). Perform Univariate and Bivariate Analysis.	3
1.2 Impute null values if present, also check for the values which are equal to zero. Do they have any meaning or do we need to change them or drop them? Check for the possibility of combining the sub levels of a ordinal variables and take actions accordingly. Explain why you are combining these sub levels with appropriate reasoning.	13
1.3 Encode the data (having string values) for Modelling. Split the data into train and test (70:30). Apply Linear regression using scikit learn. Perform checks for significant variables using appropriate method from statsmodel. Create multiple models and check the performance of Predictions on Train and Test sets using Rsquare, RMSE & Adj Rsquare. Compare these models and select the best one with appropriate reasoning.	17
1.4 Inference: Based on these predictions, what are the business insights and recommendations.	24
<b>Problem 2: Logistic Regression and LDA</b>	<b>24</b>
2.1 Data Ingestion: Read the dataset. Do the descriptive statistics and do null value condition check, write an inference on it. Perform Univariate and Bivariate Analysis. Do exploratory data analysis.	24
2.2 Do not scale the data. Encode the data (having string values) for Modelling. Data Split: Split the data into train and test (70:30). Apply Logistic Regression and LDA (linear discriminant analysis).	34
2.3 Performance Metrics: Check the performance of Predictions on Train and Test sets using Accuracy, Confusion Matrix, Plot ROC curve and get ROC_AUC score for each model Final Model: Compare Both the models and write inference which model is best/optimized.	37
2.4 Inference: Based on these predictions, what are the insights and recommendations.	44

## List of Tables :

Table-1 Dataset Sample.....	3
Table-2 Describe the data.....	5
Table-3 Data .....	15
Table-4 Scaling data .....	16
Table-5 Sample Dataset.....	25
Table-6 Describe data.....	26
Table-7 Datasets.....	34
Table-8 Datasets.....	43

## List of Figures :

Fig.1 Distplot.....	7
Fig.2 Boxplot.....	8
Fig.3 Countplot.....	10
Fig.4 Bargraph.....	11
Fig.5 Pairplot.....	12
Fig.6 Heatmap.....	13
Fig.7 Distplot.....	27
Fig.8 Boxplot.....	28
Fig.9 Countplot.....	29
Fig10. Bar Graph .....	29
Fig. 11 Distplott.....	30
Fig.12 Catplot.....	31
Fig.13 Catplot.....	32
Fig.14 Pairplot .....	33
Fig.15 Heatmap.....	34
Fig. 16 roc curve .....	40
Fig. 17 roc curve .....	41
Fig. 18 roc curve .....	43

## Problem Statement - 1 Linear Regression

You are hired by a company Gem Stones co Ltd, which is a cubic zirconia manufacturer. You are provided with the dataset containing the prices and other attributes of almost 27,000 cubic zirconia (which is an inexpensive diamond alternative with many of the same qualities as a diamond). The company is earning different profits on different prize slots. You have to help the company in predicting the price for the stone on the basis of the details given in the dataset so it can distinguish between higher profitable stones and lower profitable stones so as to have better profit share. Also, provide them with the best 5 attributes that are most important.

**1.1. Read the data and do exploratory data analysis. Describe the data briefly. (Check the null values, Data types, shape, EDA, duplicate values). Perform Univariate and Bivariate Analysis.**

Sample of the dataset

	Unnamed: 0	carat	cut	color	clarity	depth	table	x	y	z	price
0	1	0.30	Ideal	E	SI1	62.1	58.0	4.27	4.29	2.66	499
1	2	0.33	Premium	G	IF	60.8	58.0	4.42	4.46	2.70	984
2	3	0.90	Very Good	E	VVS2	62.2	60.0	6.04	6.12	3.78	6289
3	4	0.42	Ideal	F	VS1	61.6	56.0	4.82	4.80	2.96	1082
4	5	0.31	Ideal	F	VVS1	60.4	59.0	4.35	4.43	2.65	779

Table-01Dataset Sample

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 26967 entries, 0 to 26966
Data columns (total 10 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   carat       26967 non-null  float64
 1   cut         26967 non-null  object
 2   color       26967 non-null  object
 3   clarity     26967 non-null  object
 4   depth       26270 non-null  float64
 5   table       26967 non-null  float64
 6   x           26967 non-null  float64
 7   y           26967 non-null  float64
 8   z           26967 non-null  float64
 9   price       26967 non-null  int64
dtypes: float64(6), int64(1), object(3)
memory usage: 2.1+ MB

```

- From the above results we can see that there is some missing value present in the dataset.
- There are a total of 26967 rows and 10 columns in the dataset.
- 6 variables out of 10 is float64 , 1 variable is int64 and 3 variables out 10 are object data types.

#### missing values :

```

carat      0
cut         0
color      0
clarity     0
depth      697
table       0
x           0
y           0
z           0
price       0
dtype: int64

```

From the above result we can see that there are 697 values missing in the depth variable.

### Duplicate values :

```
: dups = df.duplicated()
print('Number of duplicate rows = %d' % (dups.sum()))
```

executed in 57ms, finished 14:14:17 2021-11-27

Number of duplicate rows = 0

There is no duplicate value in the datasets.

### Describe datasets :

	carat	cut	color	clarity	depth	table	x	y	z	price
<b>count</b>	26967.000000	26967	26967	26967	26270.000000	26967.000000	26967.000000	26967.000000	26967.000000	26967.000000
<b>unique</b>	NaN	5	7	8	NaN	NaN	NaN	NaN	NaN	NaN
<b>top</b>	NaN	Ideal	G	SI1	NaN	NaN	NaN	NaN	NaN	NaN
<b>freq</b>	NaN	10816	5661	6571	NaN	NaN	NaN	NaN	NaN	NaN
<b>mean</b>	0.798375	NaN	NaN	NaN	61.745147	57.456080	5.729854	5.733569	3.538057	3939.518115
<b>std</b>	0.477745	NaN	NaN	NaN	1.412860	2.232068	1.128516	1.166058	0.720624	4024.864666
<b>min</b>	0.200000	NaN	NaN	NaN	50.800000	49.000000	0.000000	0.000000	0.000000	326.000000
<b>25%</b>	0.400000	NaN	NaN	NaN	61.000000	56.000000	4.710000	4.710000	2.900000	945.000000
<b>50%</b>	0.700000	NaN	NaN	NaN	61.800000	57.000000	5.690000	5.710000	3.520000	2375.000000
<b>75%</b>	1.050000	NaN	NaN	NaN	62.500000	59.000000	6.550000	6.540000	4.040000	5360.000000
<b>max</b>	4.500000	NaN	NaN	NaN	73.600000	79.000000	10.230000	58.900000	31.800000	18818.000000

Table-02 Describe the data

```
CUT : 5
Fair      781
Good     2441
Very Good 6030
Premium   6899
Ideal    10816
Name: cut, dtype: int64
```

```
COLOR : 7
J      1443
I      2771
D      3344
H      4102
F      4729
E      4917
G      5661
Name: color, dtype: int64
```

```
CLARITY : 8
I1       365
IF       894
VVS1     1839
VVS2     2531
VS1      4093
SI2      4575
VS2      6099
SI1      6571
Name: clarity, dtype: int64
```

- We have 3 object variables and the remaining are numerical variables.
- In the cut variable there are 5 unique values, most of the data is Ideal and their frequency is 10816.
- In the color variable there are 7 unique values, most of the color is G and their frequency is 5661.
- In the clarity variable there are 8 unique values, most of the data is SI1 and their frequency is 6571.

- Some of the numerical variables have mean and median are almost equal and some of the variables have skewness.

### Univariate / Bivariate analysis :

#### Distplot :

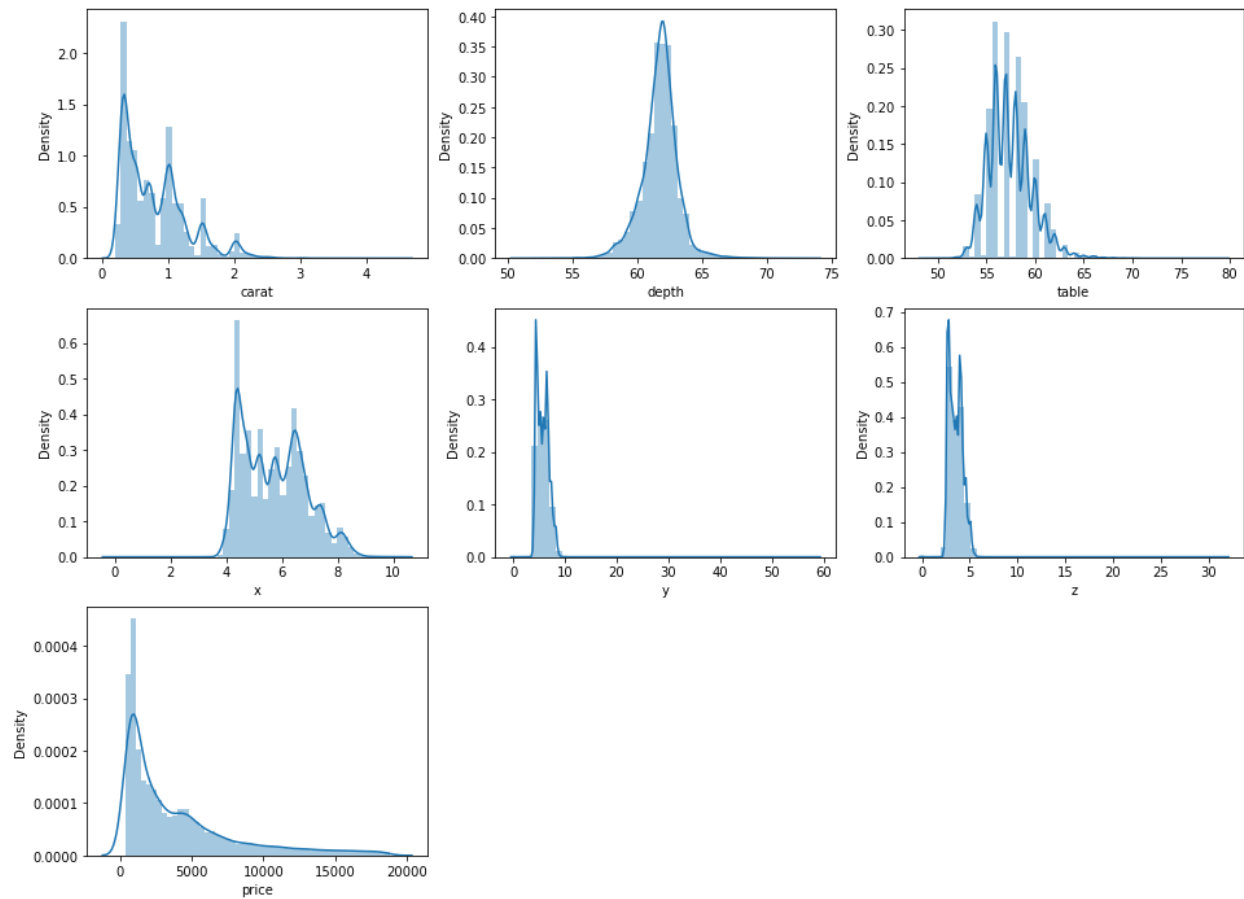


Figure-01



### Boxplot :

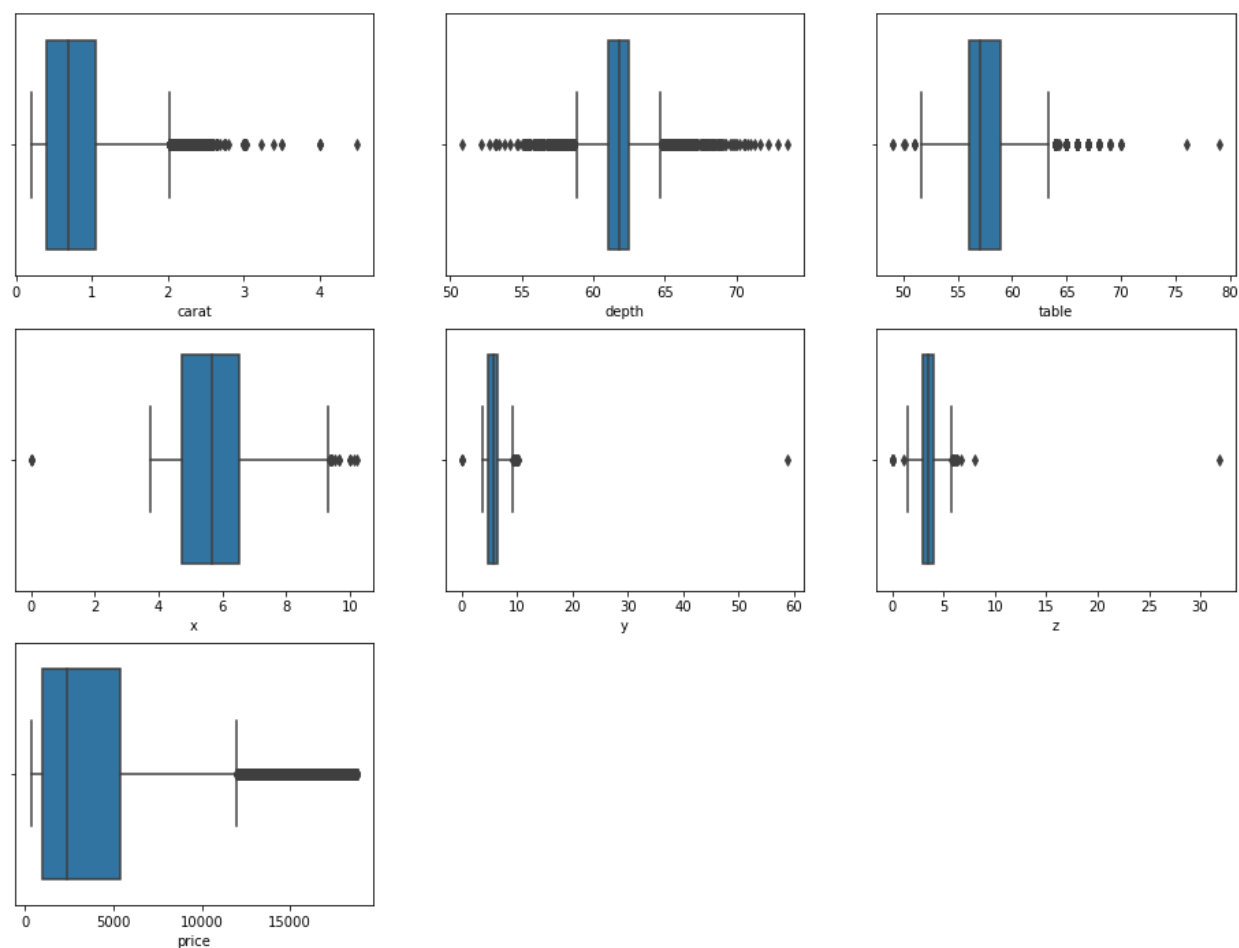


Figure-02

```
carat Skewed: 1.1165
depth Skewed: -0.0286
table Skewed: 0.7658
x Skewed: 0.388
y Skewed: 3.8502
z Skewed: 2.5683
price Skewed: 1.6185
```

### Carat feature :

- Distribution of carat variable seems to right skewed = 1.1165
- In the distplot it seems there are multiple peak points.
- There are large no. of outliers.

Depth feature :

- Distribution of depth variable seems to be left skewed = -0.0289 it is very less no. so we can say depth data is normally distributed.
- There are large no. of outliers in both the direction left and right.

Table feature :

- Distribution of table variable seems to be positively skewed =0.7658
- Boxplot of the table has outliers.

X (Length of the cubic zirconia in mm) feature :

- Distribution of the X (Length of the cubic zirconia in mm) seems positively skewed= 0.388.
- There are outliers.

Y (width of the cubic zirconia in mm) feature:

- Distribution of the Y (width of the cubic zirconia in mm) seems positively skewed= 3.8502.
- There are outliers.

Z (Height of the cubic zirconia in mm) feature :

- Distribution of the Z (Height of the cubic zirconia in mm) seems positively skewed= 2.5683.
- There are outliers.

Price feature :

- Distribution of the price feature seems positively skewed= 1.6185.
- There are outliers.

Countplot :

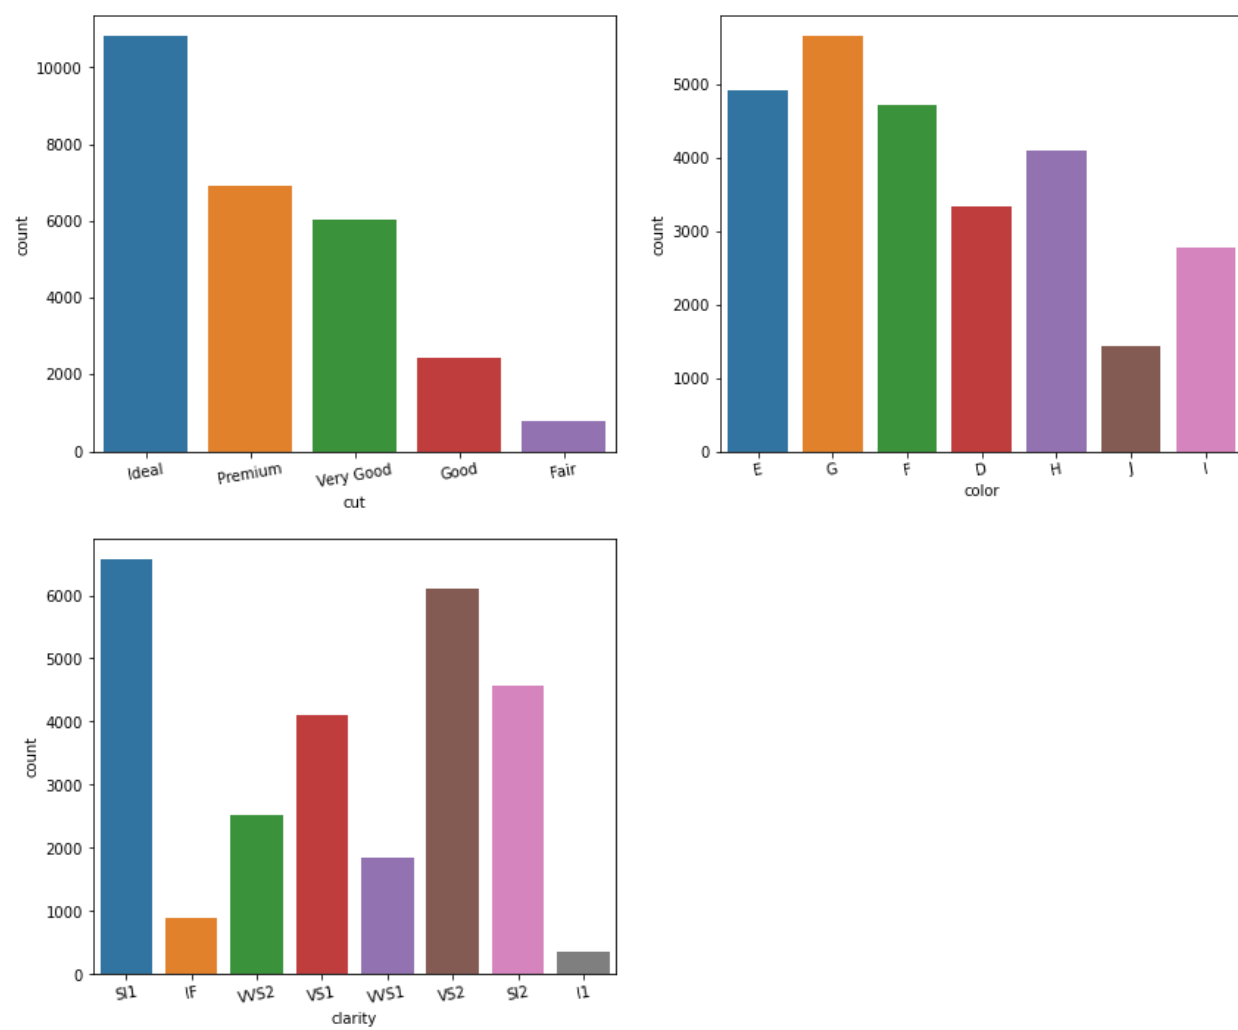


Figure-03

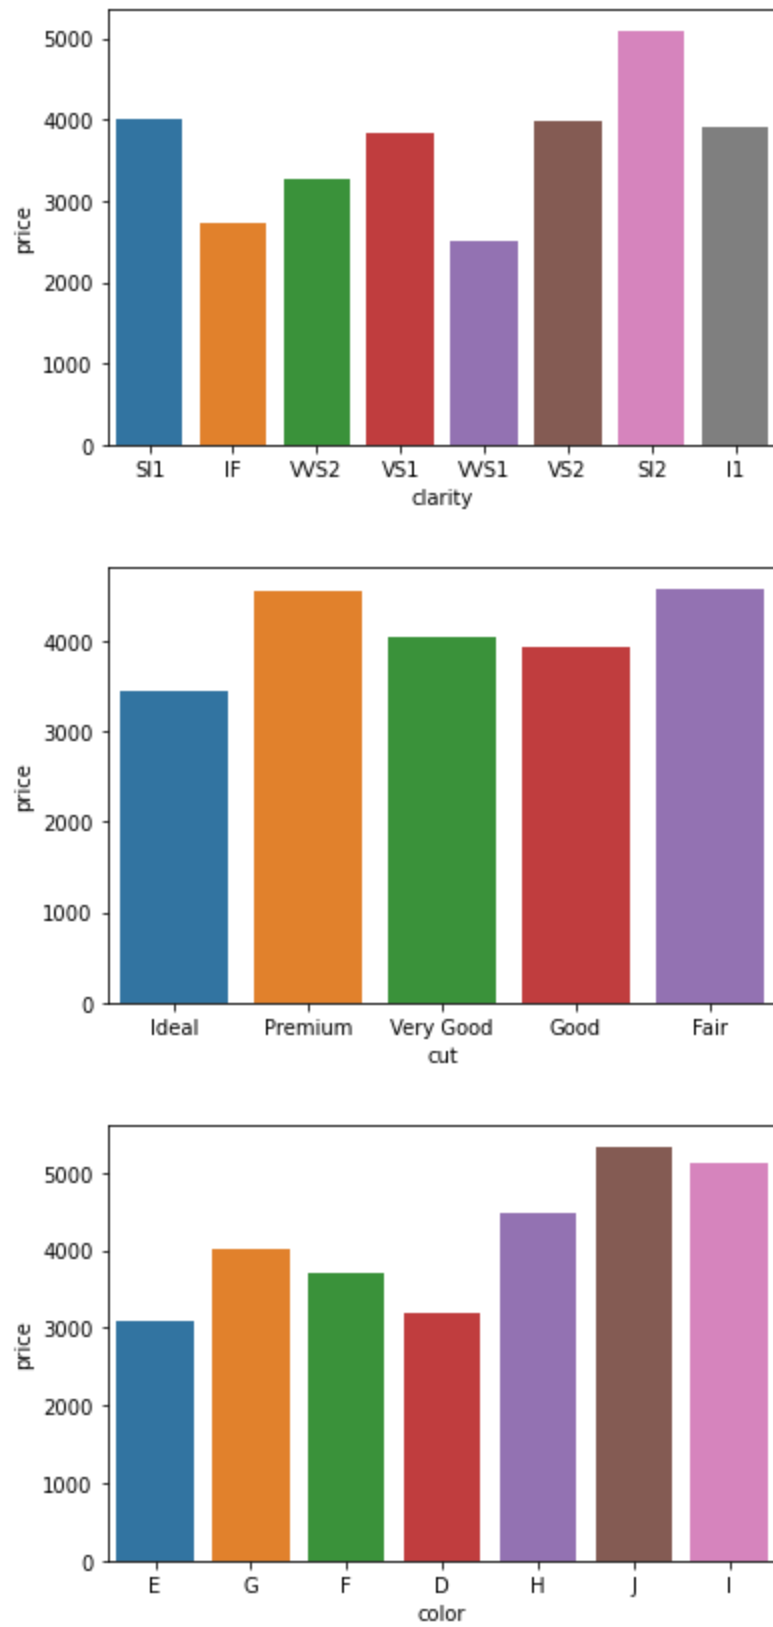


Figure-04

Cut feature :

- Ideal cut is the most preferred cut for the diamond.
- Ideal cut price is less compared to others.

Color feature :

- G color seems to be the most preferred color.
- J color price is the most expensive.

Clarity feature :

- SI1 seems to be the most preferred clarity.

Pairplot :

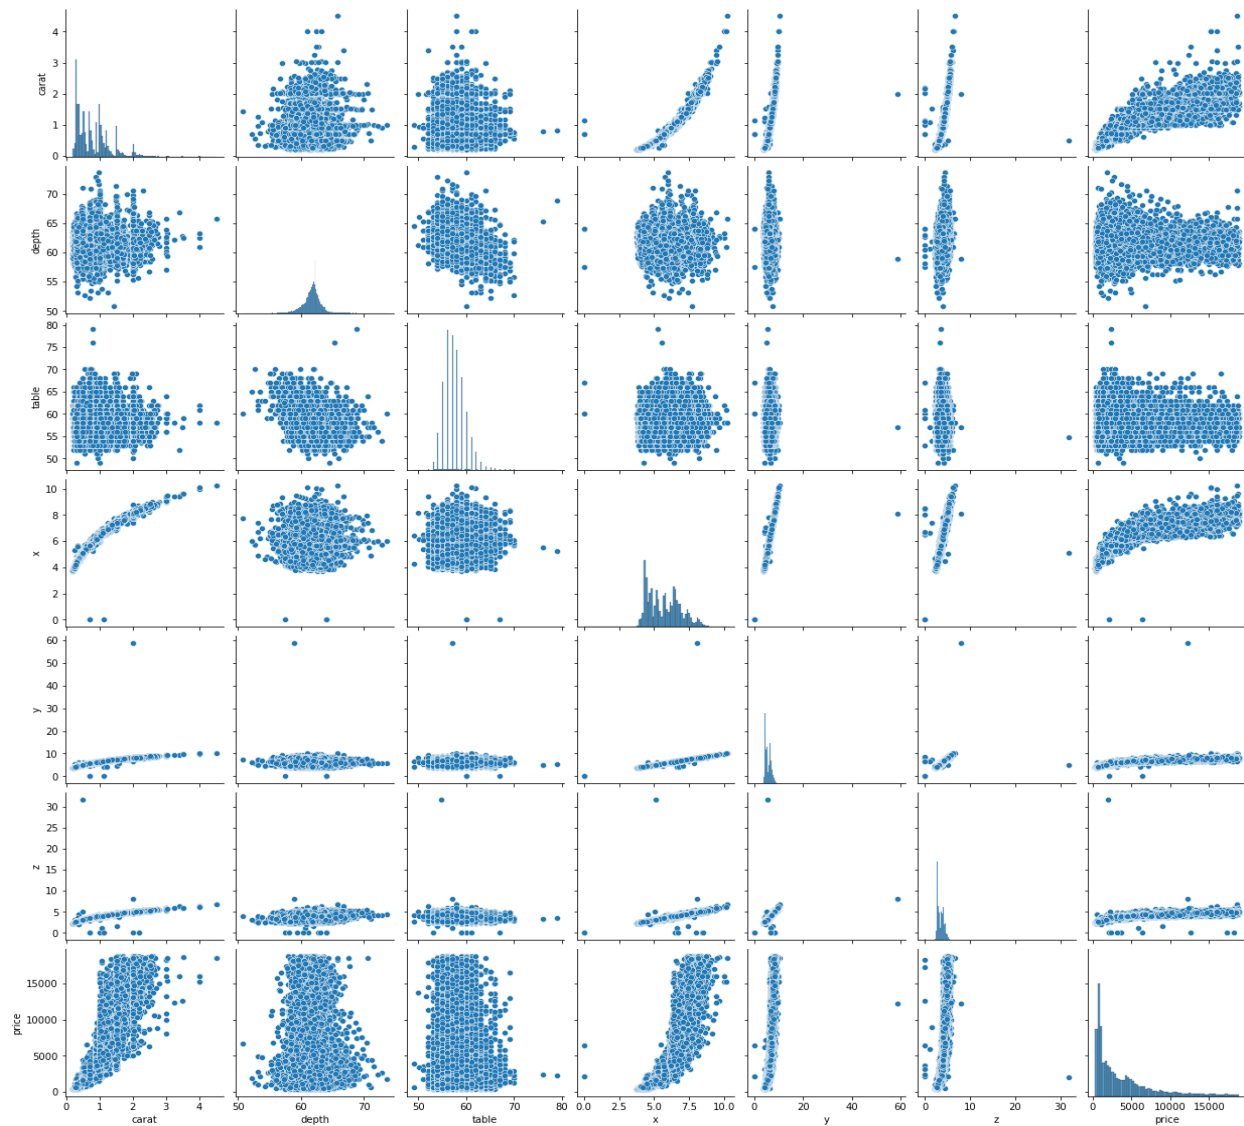


Figure -05

Heatmap :




Figure-06

After seeing the pairplot and heatmap we can say that there is presence of multicollinearity in the datasets.

**1.2 Impute null values if present, also check for the values which are equal to zero. Do they have any meaning or do we need to change them or drop them? Check for the possibility of combining the sub levels of a ordinal variables and take actions accordingly. Explain why you are combining these sub levels with appropriate reasoning.**

```
carat      0
cut        0
color      0
clarity    0
depth     697
table      0
x          0
y          0
z          0
price      0
dtype: int64
```

```
carat      0.000000
cut        0.000000
color      0.000000
clarity     0.000000
depth      0.025846
table      0.000000
x          0.000000
y          0.000000
z          0.000000
price      0.000000
dtype: float64
```

There are null values in depth variables and this is less than 5% so we can fill them using median imputation.

After the median imputation we see below there are no null values now.

```
median = df['depth'].median()
df['depth'] = df['depth'].fillna(median)

df.isnull().sum()
```

executed in 28ms, finished 16:41:05 2021-11-27

```
carat      0
cut        0
color      0
clarity     0
depth      0
table      0
x          0
y          0
z          0
price      0
dtype: int64
```

Checking if there is value that is "0" :

```
df.all()
```

executed in 17ms, finished 16:41:15 2021-11-27

```
carat      True
cut        True
color      True
clarity    True
depth      True
table      True
x          False
y          False
z          False
price      True
dtype: bool
```

	carat	cut	color	clarity	depth	table	x	y	z	price
<b>5821</b>	0.71	Good	F	SI2	64.1	60.0	0.00	0.00	0.0	2130
<b>6034</b>	2.02	Premium	H	VS2	62.7	53.0	8.02	7.95	0.0	18207
<b>6215</b>	0.71	Good	F	SI2	64.1	60.0	0.00	0.00	0.0	2130
<b>10827</b>	2.20	Premium	H	SI1	61.2	59.0	8.42	8.37	0.0	17265
<b>12498</b>	2.18	Premium	H	SI2	59.4	61.0	8.49	8.45	0.0	12631
<b>12689</b>	1.10	Premium	G	SI2	63.0	59.0	6.50	6.47	0.0	3696
<b>17506</b>	1.14	Fair	G	VS1	57.5	67.0	0.00	0.00	0.0	6381
<b>18194</b>	1.01	Premium	H	I1	58.1	59.0	6.66	6.60	0.0	3167
<b>23758</b>	1.12	Premium	G	I1	60.4	59.0	6.71	6.67	0.0	2383

Table-03

- I found that there is some row containing zero that has an X,Y,Z feature and has 0 values.
- We have to drop the zero value containing row.

Remove outliers :

- As we know there are outliers in the datasets so we have to impute the outliers.



- After the imputation of outliers we have seen there is now outliers now.




Figure-07

Scaling :

- Scaling is required for the linear regression because it reduces the multicollinearity.
- I use StandardScaler for scaling the datasets.

	carat	cut	color	clarity	depth	table	x	y	z	price
0	-1.067306	Ideal	E	SI1	0.286726	0.261941	-1.296438	-1.289580	-1.261448	-0.933219
1	-1.002414	Premium	G	IF	-0.780109	0.261941	-1.163237	-1.137532	-1.203982	-0.793428
2	0.230546	Very Good	E	VVS2	0.368790	1.189304	0.275339	0.347170	0.347606	0.735631
3	-0.807736	Ideal	F	VS1	-0.123596	-0.665422	-0.808033	-0.833436	-0.830451	-0.765181
4	-1.045675	Ideal	F	VVS1	-1.108366	0.725622	-1.225398	-1.164364	-1.275814	-0.852515

Table-04

**1.3 Encode the data (having string values) for Modelling. Split the data into train and test (70:30). Apply Linear regression using scikit learn. Perform checks for significant variables using appropriate method from statsmodel. Create multiple models and check the performance of Predictions on Train and Test sets using Rsquare, RMSE & Adj Rsquare. Compare these models and select the best one with appropriate reasoning.**

Encode the string variables :

After encoding the datasets then shape is 5X24

```
data.columns
```

executed in 16ms, finished 18:59:41 2021-11-27

```
Index(['carat', 'depth', 'table', 'x', 'y', 'z', 'price', 'cut_Good',
      'cut_Ideal', 'cut_Premium', 'cut_Very Good', 'color_E', 'color_F',
      'color_G', 'color_H', 'color_I', 'color_J', 'clarity_IF', 'clarity_SI1',
      'clarity_SI2', 'clarity_VS1', 'clarity_VS2', 'clarity_VS1',
      'clarity_VS2'],
      dtype='object')
```

Dropping target variables for the model and split the data 70:30 ratio for the train and test respectively.

```
X = data.drop('price', axis=1)

# Copy target into the y dataframe.
y = data[['price']]
```

executed in 22ms, finished 19:02:53 2021-11-27

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3 , random_state=1)
```

executed in 26ms, finished 19:02:55 2021-11-27

Linear Regression Model :

```
: regression_model = LinearRegression()  
  regression_model.fit(X_train, y_train)
```

executed in 50ms, finished 19:02:55 2021-11-27

```
: LinearRegression()
```

Coefficient of variables :

The coefficient for carat is 1.2358093063484565  
The coefficient for depth is 0.005681781619185346  
The coefficient for table is -0.0149297410250086  
The coefficient for x is -0.35355573073496177  
The coefficient for y is 0.338428967376515  
The coefficient for z is -0.15648799801269386  
The coefficient for cut\_Good is 0.10904474596752811  
The coefficient for cut\_Ideal is 0.1766242155294696  
The coefficient for cut\_Premium is 0.1722373351767154  
The coefficient for cut\_Very Good is 0.14592655610103092  
The coefficient for color\_E is -0.054565752180473735  
The coefficient for color\_F is -0.072690724250613  
The coefficient for color\_G is -0.1167998797001075  
The coefficient for color\_H is -0.24082413839492464  
The coefficient for color\_I is -0.37566722366410227  
The coefficient for color\_J is -0.5433905358981461  
The coefficient for clarity\_IF is 1.1593629758064512  
The coefficient for clarity\_SI1 is 0.7409791729531008  
The coefficient for clarity\_SI2 is 0.49817342838278567  
The coefficient for clarity\_VS1 is 0.9718720807450599  
The coefficient for clarity\_VS2 is 0.8883054566346622  
The coefficient for clarity\_VVS1 is 1.092461357490716  
The coefficient for clarity\_VVS2 is 1.0800417891526692

Intercept of model :

```
intercept = regression_model.intercept_[0]  
  
print("The intercept for our model is {}".format(intercept))
```

executed in 23ms, finished 19:02:57 2021-11-27

The intercept for our model is -0.8323524295093438

```
: # R square on training data  
regression_model.score(X_train, y_train)
```

executed in 36ms, finished 19:02:57 2021-11-27

```
: 0.9419557931252712
```

```
# R square on testing data  
regression_model.score(X_test, y_test)
```

executed in 40ms, finished 19:02:58 2021-11-27

0.9381643998102491

```
#RMSE on Training data  
predicted_train=regression_model.fit(X_train, y_train).predict(X_train)  
np.sqrt(metrics.mean_squared_error(y_train,predicted_train))
```

executed in 42ms, finished 19:02:59 2021-11-27

0.23992842984922294

```
#RMSE on Testing data  
predicted_test=regression_model.fit(X_train, y_train).predict(X_test)  
np.sqrt(metrics.mean_squared_error(y_test,predicted_test))
```

executed in 75ms, finished 19:02:59 2021-11-27

0.25103473833743095

VIF (variance inflation factor) :

```
carat ---> 33.35287649550623
depth ---> 4.574003842337532
table ---> 1.7722022611198982
x ---> 463.94494858728734
y ---> 463.08309600508517
z ---> 238.6002431605187
cut_Good ---> 3.6104961328079184
cut_Ideal ---> 14.347409690217939
cut_Premium ---> 8.623207030351887
cut_Very Good ---> 7.852218650260111
color_E ---> 2.3710537954581707
```

We can see in VIF there is multicollinearity in the datasets from the state models. We can understand which feature is not more important for the model. We can remove those variables . Ideal value of VIF is less than 10%.

Using Statsmodel library :

```
Intercept      -0.832352
carat          1.235809
depth          0.005682
table         -0.014930
x             -0.353556
y              0.338429
z             -0.156488
cut_Good       0.109045
cut_Ideal      0.176624
cut_Premium    0.172237
cut_Very_Good  0.145927
color_E       -0.054566
color_F       -0.072691
color_G       -0.116800
color_H       -0.240824
color_I       -0.375667
color_J       -0.543391
clarity_IF     1.159363
clarity_SI1    0.740979
clarity_SI2    0.498173
clarity_VS1    0.971872
clarity_VS2    0.888305
clarity_VVS1   1.092461
clarity_VVS2   1.080042
dtype: float64
```

OLS Regression Results						
=====						
Dep. Variable:	price	R-squared:	0.942			
Model:	OLS	Adj. R-squared:	0.942			
Method:	Least Squares	F-statistic:	1.330e+04			
Date:	Sat, 27 Nov 2021	Prob (F-statistic):	0.00			
Time:	19:40:14	Log-Likelihood:	159.94			
No. Observations:	18870	AIC:	-271.9			
Df Residuals:	18846	BIC:	-83.60			
Df Model:	23					
Covariance Type:	nonrobust					
=====						
	coef	std err	t	P> t	[0.025	0.975]
-----						
Intercept	-0.8324	0.019	-44.588	0.000	-0.869	-0.796
carat	1.2358	0.010	121.892	0.000	1.216	1.256
depth	0.0057	0.004	1.525	0.127	-0.002	0.013
table	-0.0149	0.002	-6.356	0.000	-0.020	-0.010
x	-0.3536	0.037	-9.531	0.000	-0.426	-0.281
y	0.3384	0.038	8.934	0.000	0.264	0.413
z	-0.1565	0.027	-5.742	0.000	-0.210	-0.103
cut_Good	0.1090	0.012	8.755	0.000	0.085	0.133
cut_Ideal	0.1766	0.012	14.581	0.000	0.153	0.200
cut_Premium	0.1722	0.012	14.785	0.000	0.149	0.195
cut_Very_Good	0.1459	0.012	12.269	0.000	0.123	0.169
color_E	-0.0546	0.006	-8.429	0.000	-0.067	-0.042
color_F	-0.0727	0.007	-11.075	0.000	-0.086	-0.060
color_G	-0.1168	0.006	-18.258	0.000	-0.129	-0.104
color_H	-0.2408	0.007	-35.323	0.000	-0.254	-0.227
color_I	-0.3757	0.008	-49.521	0.000	-0.391	-0.361
color_J	-0.5434	0.009	-58.186	0.000	-0.562	-0.525
clarity_IF	1.1594	0.019	62.524	0.000	1.123	1.196
clarity_SI1	0.7410	0.016	46.643	0.000	0.710	0.772
clarity_SI2	0.4982	0.016	31.177	0.000	0.467	0.529
clarity_VS1	0.9719	0.016	59.986	0.000	0.940	1.004
clarity_VS2	0.8883	0.016	55.618	0.000	0.857	0.920
clarity_VVS1	1.0925	0.017	63.630	0.000	1.059	1.126
clarity_VVS2	1.0800	0.017	64.730	0.000	1.047	1.113
=====						
Omnibus:	4696.785	Durbin-Watson:	1.994			
Prob(Omnibus):	0.000	Jarque-Bera (JB):	17654.853			
Skew:	1.208	Prob(JB):	0.00			
Kurtosis:	7.076	Cond. No.	57.3			

In this model we found that depth  $p > |t|$  is greater than 0.5 so we should remove and again build the model.

```

Intercept      -0.832269
carat          1.236946
table          -0.015612
x              -0.365810
y              0.315607
z              -0.122330
cut_Good        0.110307
cut_Ideal       0.175305
cut_Premium     0.170878
cut_Very_Good   0.145550
color_E         -0.054635
color_F         -0.072740
color_G         -0.116725
color_H         -0.240702
color_I         -0.375359
color_J         -0.543203
clarity_IF      1.159684
clarity_SI1     0.741936
clarity_SI2     0.498868
clarity_VS1     0.972497
clarity_VS2     0.889069
clarity_VVS1    1.092825
clarity_VVS2    1.080657
dtype: float64

```

```
print(lm2.summary())
```

executed in 41ms, finished 19:56:43 2021-11-27

#### OLS Regression Results

```

=====
Dep. Variable:          price    R-squared:                0.942
Model:                  OLS      Adj. R-squared:           0.942
Method:                 Least Squares    F-statistic:            1.390e+04
Date:                   Sat, 27 Nov 2021    Prob (F-statistic):      0.00
Time:                   19:56:43    Log-Likelihood:         158.78
No. Observations:      18870    AIC:                    -271.6
Df Residuals:          18847    BIC:                    -91.12
Df Model:               22
Covariance Type:       nonrobust

```

	coef	std err	t	P> t	[0.025	0.975]
Intercept	-0.8323	0.019	-44.583	0.000	-0.869	-0.796
carat	1.2369	0.010	122.331	0.000	1.217	1.257
table	-0.0156	0.002	-6.770	0.000	-0.020	-0.011
x	-0.3658	0.036	-10.101	0.000	-0.437	-0.295
y	0.3156	0.035	9.069	0.000	0.247	0.384
z	-0.1223	0.016	-7.883	0.000	-0.153	-0.092
cut_Good	0.1103	0.012	8.876	0.000	0.086	0.135
cut_Ideal	0.1753	0.012	14.508	0.000	0.152	0.199
cut_Premium	0.1709	0.012	14.711	0.000	0.148	0.194
cut_Very_Good	0.1455	0.012	12.239	0.000	0.122	0.169
color_E	-0.0546	0.006	-8.439	0.000	-0.067	-0.042
color_F	-0.0727	0.007	-11.082	0.000	-0.086	-0.060
color_G	-0.1167	0.006	-18.246	0.000	-0.129	-0.104
color_H	-0.2407	0.007	-35.306	0.000	-0.254	-0.227
color_I	-0.3754	0.008	-49.497	0.000	-0.390	-0.360
color_J	-0.5432	0.009	-58.169	0.000	-0.562	-0.525
clarity_IF	1.1597	0.019	62.544	0.000	1.123	1.196
clarity_SI1	0.7419	0.016	46.738	0.000	0.711	0.773
clarity_SI2	0.4989	0.016	31.232	0.000	0.468	0.530
clarity_VS1	0.9725	0.016	60.042	0.000	0.941	1.004
clarity_VS2	0.8891	0.016	55.691	0.000	0.858	0.920
clarity_VVS1	1.0928	0.017	63.655	0.000	1.059	1.126
clarity_VVS2	1.0807	0.017	64.784	0.000	1.048	1.113
=====						
Omnibus:		4699.504	Durbin-Watson:			1.994
Prob(Omnibus):		0.000	Jarque-Bera (JB):			17704.272
Skew:		1.208	Prob(JB):			0.00
Kurtosis:		7.084	Cond. No.			56.9

Conclusion :

The final Linear Regression equation is:

Log\_price = (-0.83) \* Intercept + (1.24) \* carat + (-0.02) \* table + (-0.37) \* x + (0.32) \* y + (-0.12) \* z + (0.11) \* cut\_Good + (0.18) \* cut\_Ideal + (0.17) \* cut\_Premium + (0.15) \* cut\_Very\_Good + (-0.05) \* color\_E + (-0.07) \* color\_F + (-0.12) \* color\_G + (-0.24) \* color\_H + (-0.38) \* color\_I + (-0.54) \* color\_J + (1.16) \* clarity\_IF + (0.74) \* clarity\_SI1 + (0.5) \* clarity\_SI2 + (0.97) \* clarity\_VS1 + (0.89) \* clarity\_VS2 + (1.09) \* clarity\_VVS1 + (1.08) \* clarity\_VVS2

Remarks : Dropping variables would bring down the multicollinearity level down.



## 1.4 Inference: Based on these predictions, what are the business insights and recommendations.

**Please explain and summarise the various steps performed in this project. There should be proper business interpretation and actionable insights present.**

- From the EDA analysis we understand that ideal cuts have more profit to the company.
- Color H, I, J have more profit for the company.
- need to run promotional marketing campaigns or evaluate if we need to tie up with an alternate agency. It will increase sales.
- Using sats model, if we could run the model again we can have P values and coefficients which will give us better understanding of the relationship, so that values more 0.05 we can drop those variables and run again the model for better result.
- The ideal, premium, very good, cut types are those which are bringing profit so that we could use marketing for these to bring us more profits.

## Problem 2: Logistic Regression and LDA

You are hired by a tour and travel agency which deals in selling holiday packages. You are provided details of 872 employees of a company. Among these employees, some opted for the package and some didn't. You have to help the company in predicting whether an employee will opt for the package or not on the basis of the information given in the data set. Also, find out the important factors on the basis of which the company will focus on particular employees to sell their packages.

### 2.1 Data Ingestion: Read the dataset. Do the descriptive statistics and do null value condition check, write an inference on it. Perform Univariate and Bivariate Analysis. Do exploratory data analysis.

### Sample of the dataset:

	Unnamed: 0	Holliday_Package	Salary	age	educ	no_young_children	no_older_children	foreign
0	1	no	48412	30	8	1	1	no
1	2	yes	37207	45	8	0	1	no
2	3	no	58022	46	9	0	0	no
3	4	no	66503	31	11	2	0	no
4	5	no	66734	44	12	0	2	no

Table-05 Sample datasets

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 872 entries, 0 to 871
Data columns (total 7 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Holliday_Package      872 non-null   object
1   Salary                872 non-null   int64
2   age                  872 non-null   int64
3   educ                 872 non-null   int64
4   no_young_children     872 non-null   int64
5   no_older_children     872 non-null   int64
6   foreign               872 non-null   object
dtypes: int64(5), object(2)
memory usage: 47.8+ KB
```

- From the above results we can see that there is no missing value present in the dataset.
- There are a total of 872 rows and 7 columns in the dataset.
- 2 variables out of 7 are objects and 5 variables are int64 data types.

### Missing values :

```
Holliday_Package    0
Salary              0
age                 0
educ                0
no_young_children   0
no_older_children   0
foreign             0
dtype: int64
```

From the above result we can see that there is no missing value in the datasets.

### Duplicate values :

```
dups = df.duplicated()
print('Number of duplicate rows = %d' % (dups.sum()))
```

executed in 17ms, finished 16:32:37 2021-11-28

Number of duplicate rows = 0

There is no duplicate value in the datasets.

### Describe datasets :

	Holliday_Package	Salary	age	educ	no_young_children	no_older_children	foreign
count	872	872.000000	872.000000	872.000000	872.000000	872.000000	872
unique	2	NaN	NaN	NaN	NaN	NaN	2
top	no	NaN	NaN	NaN	NaN	NaN	no
freq	471	NaN	NaN	NaN	NaN	NaN	656
mean	NaN	47729.172018	39.955275	9.307339	0.311927	0.982798	NaN
std	NaN	23418.668531	10.551675	3.036259	0.612870	1.086786	NaN
min	NaN	1322.000000	20.000000	1.000000	0.000000	0.000000	NaN
25%	NaN	35324.000000	32.000000	8.000000	0.000000	0.000000	NaN
50%	NaN	41903.500000	39.000000	9.000000	0.000000	1.000000	NaN
75%	NaN	53469.500000	48.000000	12.000000	0.000000	2.000000	NaN
max	NaN	236961.000000	62.000000	21.000000	3.000000	6.000000	NaN

Table-06 Describe data

HOLLIDAY\_PACKAGE : 2

yes 401

no 471

Name: Holliday\_Package, dtype: int64

FOREIGN : 2

yes 216

no 656

Name: foreign, dtype: int64

- We have 2 object variables and the remaining are numerical variables.
- The Holloday\_package variable is target variable.
- In the FOREIGN variable most of the data is non-foreigner.

### Univariate / Bivariate analysis :

Distplot :

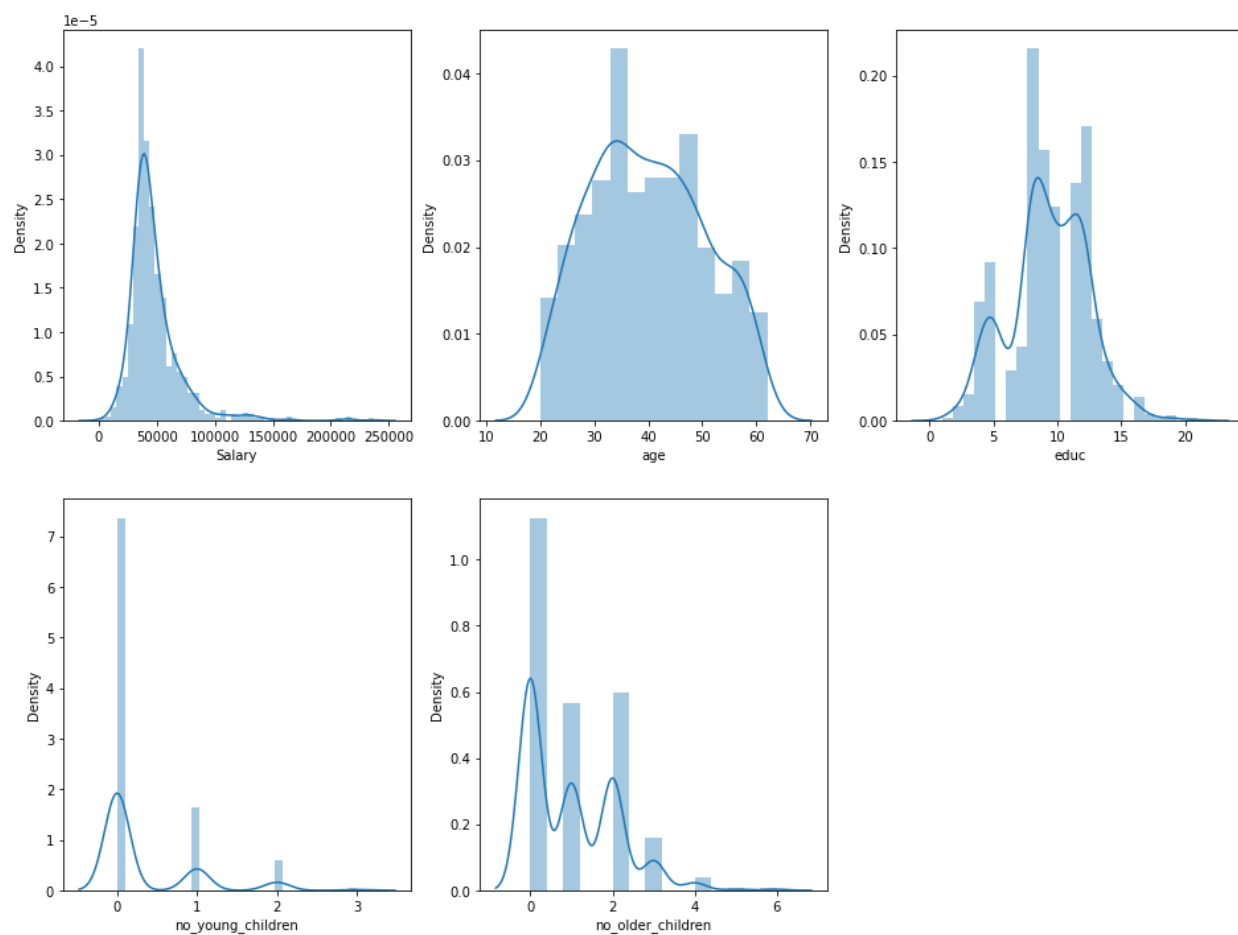


Figure-07

Boxplot :

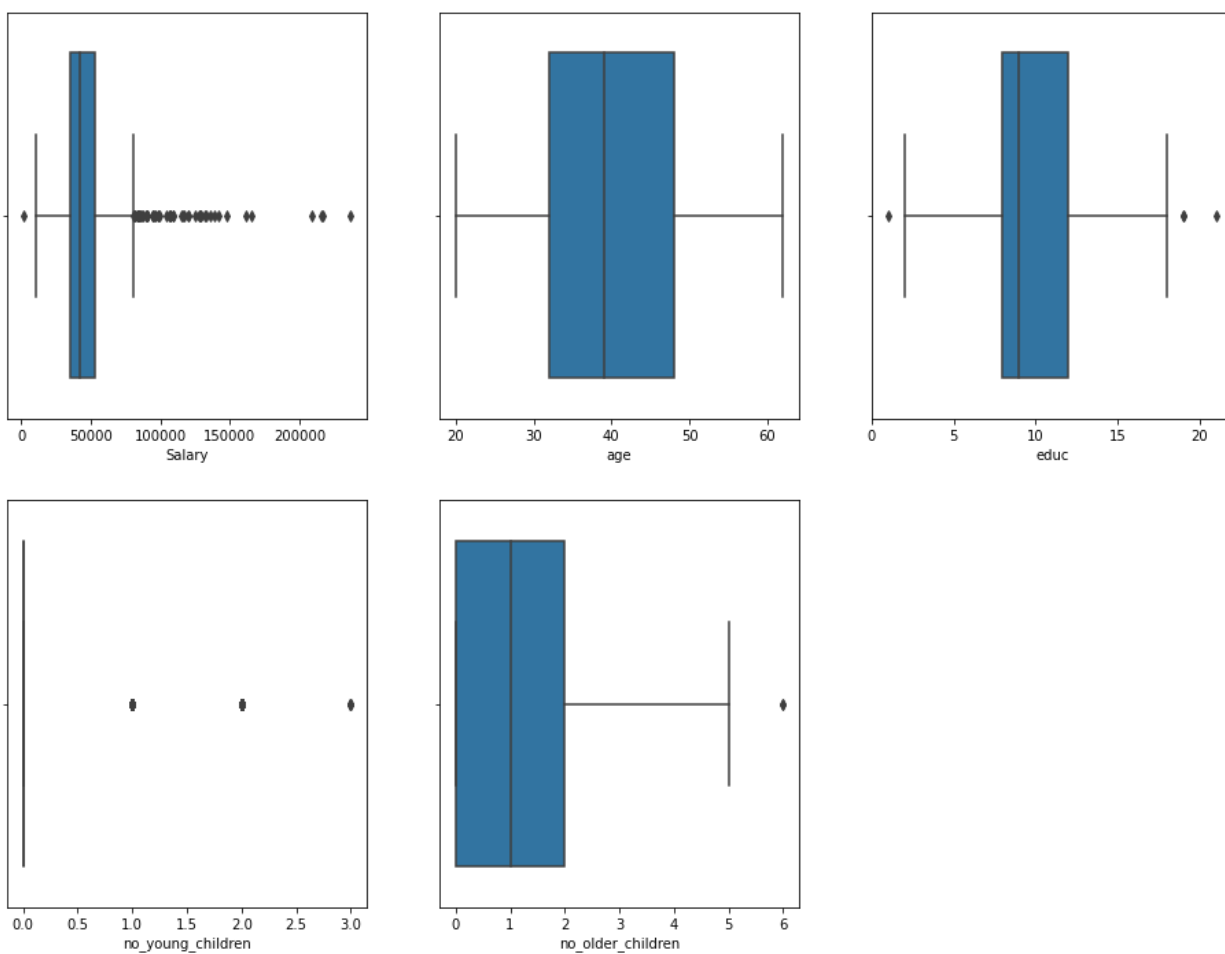


Figure-08

Salary Feature :

- Distribution of salary feature seems right skewed.
- In the distplot there seem to be outliers.

Age feature :

- Distribution of age feature seems normally distributed.
- In the distplot there seems to be no outliers.

Educ feature :

- Distribution of educ features there are multiple peak points.
- There are some outliers in the datasets.

Countplot :

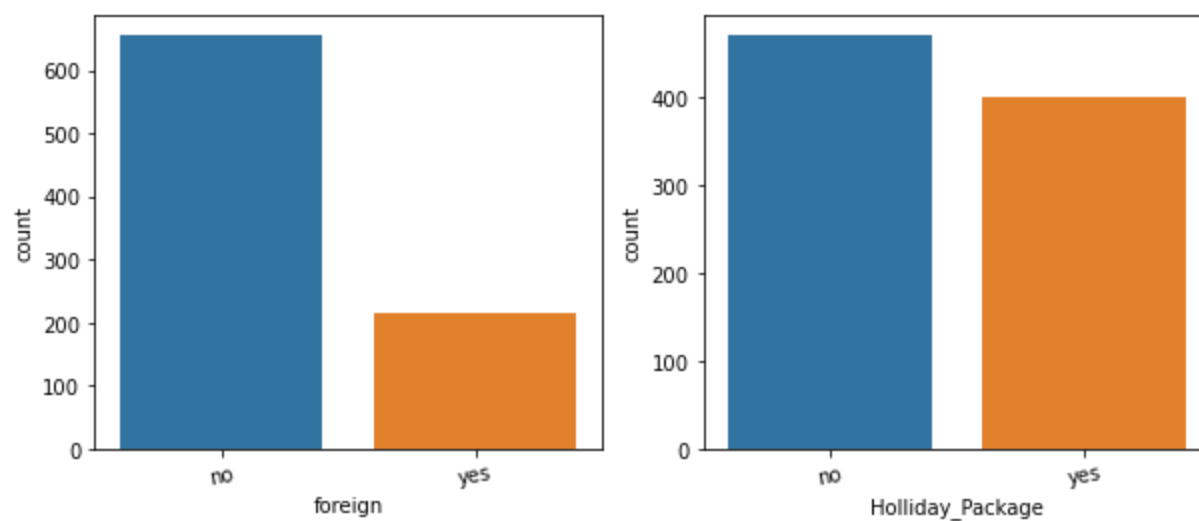


Figure-09

Bivariant :

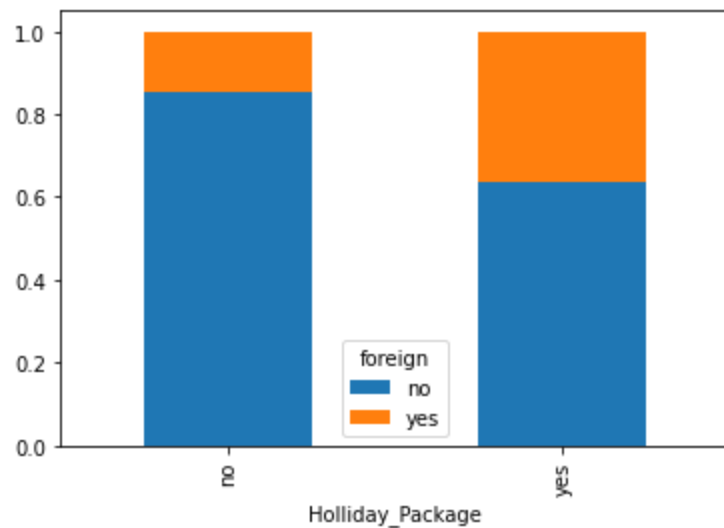


Figure-10

In the above figure we can say that foreigners are more interested in holiday packages.

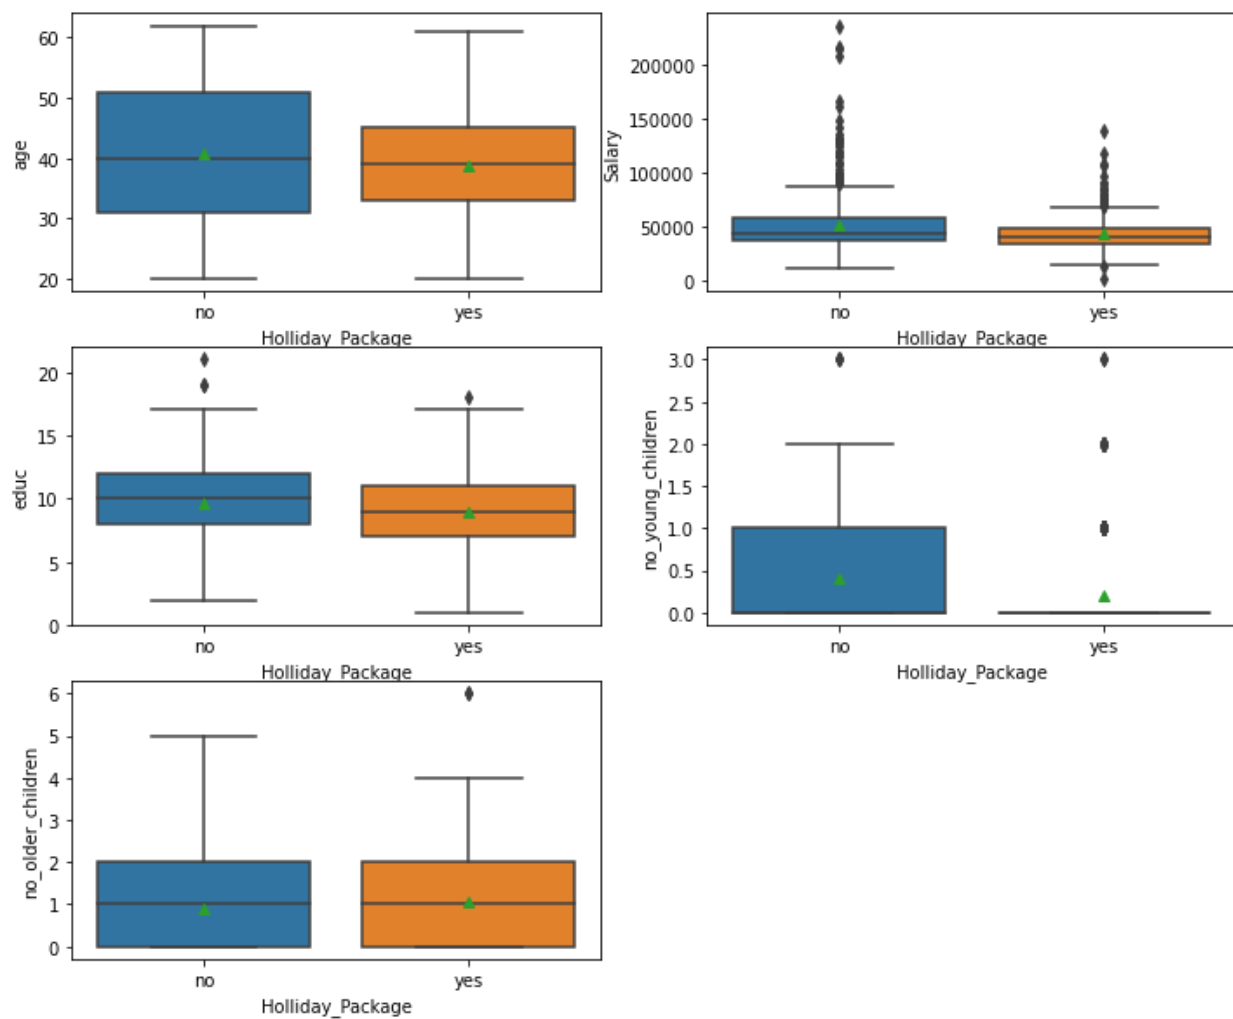


Figure -11




Figure -12






Figure-13

Pairplot :

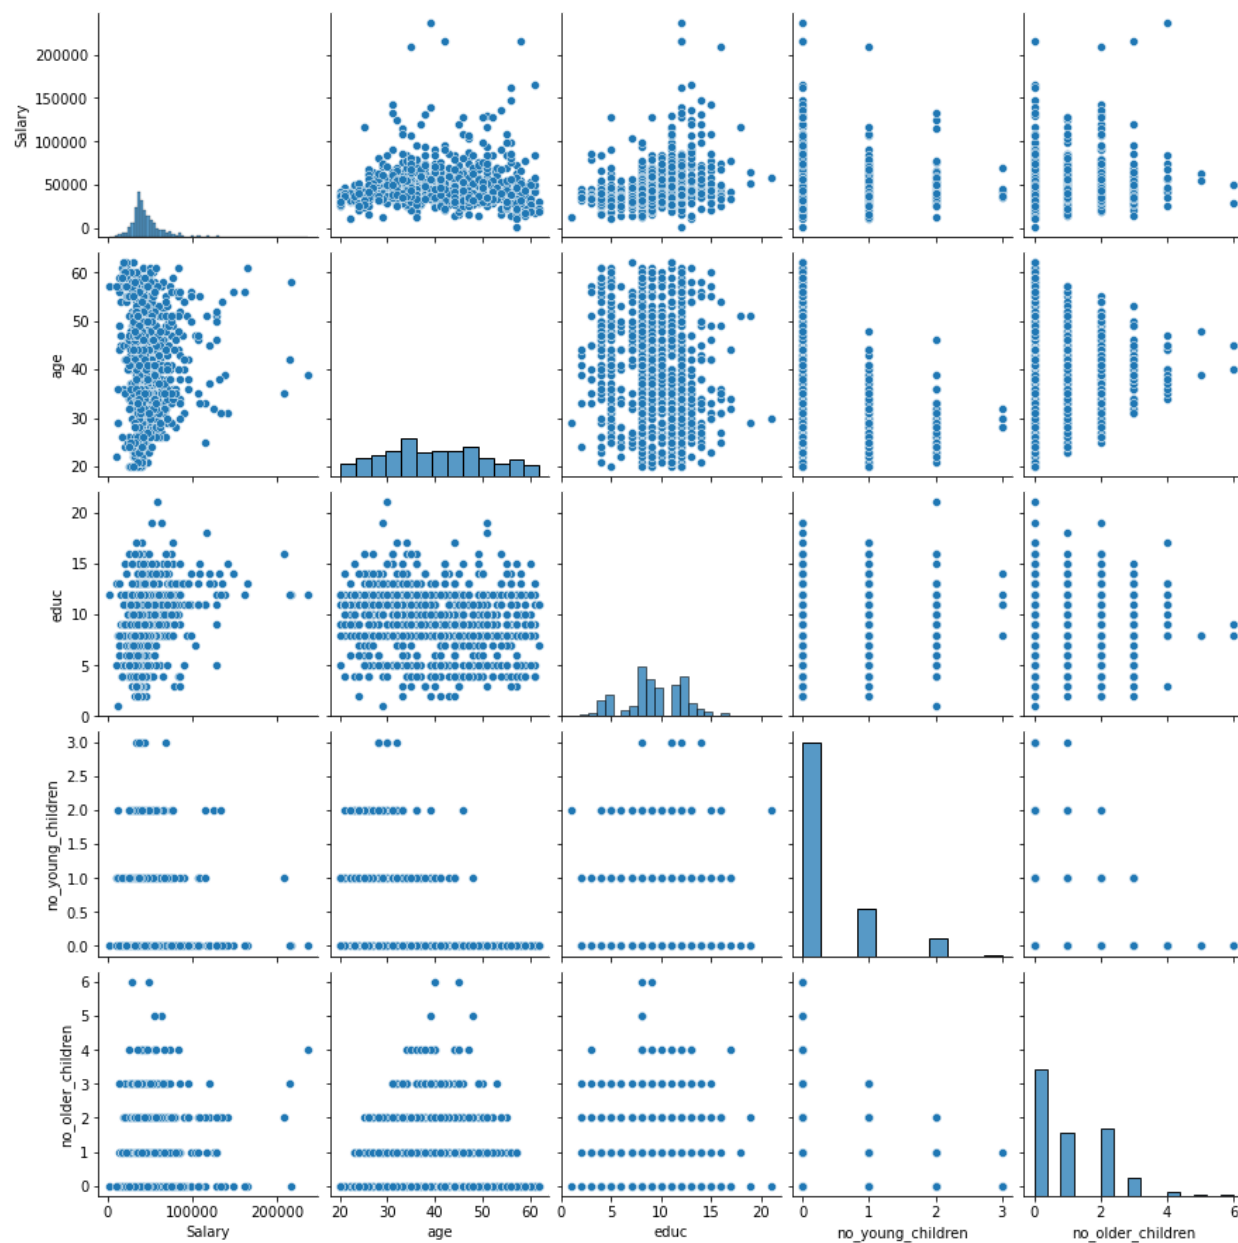


Figure - 14

Heatmap :




Figure-15

After seeing the pairplot and heatmap we can say that there is presence of multicollinearity in the datasets.

**2.2 Do not scale the data. Encode the data (having string values) for Modelling. Data Split: Split the data into train and test (70:30). Apply Logistic Regression and LDA (linear discriminant analysis).**

Encode the string variables :

	Holiday_Package	Salary	age	educ	no_young_children	no_older_children	foreign
0	0	48412.0	30	8	1	1	0
1	1	37207.0	45	8	0	1	0
2	0	58022.0	46	9	0	0	0
3	0	66503.0	31	11	2	0	0
4	0	66734.0	44	12	0	2	0

Table-07

Dropping target variables for the model and split the data 70:30 ratio for the train and test respectively.

```
: # Copy all the predictor variables into X dataframe
X = df.drop('Holliday_Package', axis=1)

# Copy target into the y dataframe.
y = df['Holliday_Package']
executed in 17ms, finished 16:33:03 2021-11-28
```

---

```
: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.30 , random_state=1,stratify=y)
executed in 12ms, finished 16:33:05 2021-11-28
```

Logistic Regression :

Grid Search Method :

Grid search method is used for logistic regression to find the optimal solving and the parameters.

```
: grid={'penalty':['l1','l2','none'],
        'solver':['lbfgs', 'liblinear'],
        'tol':[0.0001,0.000001]}
executed in 15ms, finished 16:33:05 2021-11-28
```

```
print(grid_search.best_params_,'\n')
print(grid_search.best_estimator_)
```

executed in 14ms, finished 16:33:10 2021-11-28

```
{'penalty': 'l1', 'solver': 'liblinear', 'tol': 1e-06}
```

```
LogisticRegression(max_iter=100000, n_jobs=2, penalty='l1', solver='liblinear',
                    tol=1e-06)
```

The Grid search method give, solver= liblinear, penalty= l1

Predicting the training data :

ytrain\_predict

executed in 16ms, finished 16:33:15 2021-11-28

```
array([1, 0, 0, 1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 0, 1, 0, 1, 1, 1, 0,
      1, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 1,
      1, 0, 0, 1, 0, 1, 1, 0, 1, 0, 0, 1, 0, 1, 1, 0, 1, 1, 1, 0, 1, 1,
      0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 1, 1, 1, 1, 0, 1, 1, 0, 1, 1, 0, 0,
      0, 0, 1, 1, 1, 1, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 1, 0, 0,
      1, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 0, 0, 1, 0, 1, 0,
      1, 0, 1, 1, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 1, 1, 0, 1, 1, 0, 0,
      1, 0, 0, 1, 0, 1, 0, 0, 0, 1, 1, 0, 0, 1, 1, 1, 0, 0, 1, 0, 0, 0,
      1, 0, 1, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 1, 1, 0, 1, 1, 1,
      1, 0, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 0, 0, 1, 0, 1, 1, 0, 0,
      0, 0, 0, 1, 0, 1, 0, 1, 0, 0, 1, 1, 0, 0, 1, 0, 1, 1, 1, 0, 1, 1,
      1, 1, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 0, 1, 0, 0,
      0, 1, 1, 0, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0, 0,
      0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 1,
      0, 0, 0, 1, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1, 0, 1,
      0, 0, 1, 1, 0, 0, 0, 1, 1, 1, 1, 1, 1, 0, 1, 0, 0, 0, 1, 1, 1, 0,
      1, 0, 0, 1, 0, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 1, 0,
      1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 1, 1, 0, 1, 0, 1, 0, 1, 0, 1,
      0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 0, 0, 1,
      1, 1, 1, 0, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0, 1, 0, 1, 1, 1, 0, 1, 0,
      0, 1, 0, 0, 0, 0, 0, 1, 0, 0, 1, 1, 0, 0, 1, 0, 0, 0, 0, 0, 1, 0,
      0, 0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 0, 0, 1, 1, 0], dtype=int8)
```

LDA MODEL :

```
#Build LDA Model  
clf = LinearDiscriminantAnalysis()  
model=clf.fit(X_train,y_train)
```

executed in 20ms, finished 16:33:18 2021-11-28

**2.3 Performance Metrics: Check the performance of Predictions on Train and Test sets using Accuracy, Confusion Matrix, Plot ROC curve and get ROC\_AUC score for each model Final Model: Compare Both the models and write inference which model is best/optimized.**

Logistic Regression Performance :


Confusion matrix on training data :

```
## Confusion matrix on the training data
```

```
plot_confusion_matrix(best_model,X_train,y_train)
print(classification_report(y_train, ytrain_predict),'\n');
```

executed in 163ms, finished 16:33:49 2021-11-28

	precision	recall	f1-score	support
0	0.67	0.74	0.71	329
1	0.66	0.57	0.61	281
accuracy			0.66	610
macro avg	0.66	0.66	0.66	610
weighted avg	0.66	0.66	0.66	610




Confusion matrix on testing data :

```
## Confusion matrix on the test data
```

```
plot_confusion_matrix(best_model,X_test,y_test)
print(classification_report(y_test, ytest_predict),'\n');
```

executed in 158ms, finished 16:33:49 2021-11-28

	precision	recall	f1-score	support
0	0.65	0.77	0.71	142
1	0.66	0.52	0.58	120
accuracy			0.66	262
macro avg	0.66	0.65	0.64	262
weighted avg	0.66	0.66	0.65	262





Accuracy :

```
# Accuracy - Training Data
```

```
lr_train_acc = best_model.score(X_train, y_train)  
lr_train_acc
```

```
executed in 10ms, finished 16:33:50 2021-11-28
```

0.6639344262295082

AUC , ROC curve for train data :

AUC: 0.733




Figure-16

AUC, ROC cure for test data :

AUC: 0.716




Figure- 17

LDA MODEL Performance :

Classification report on training data :

```
print(classification_report(y_train, pred_class_train))
```

executed in 16ms, finished 16:33:55 2021-11-28

	precision	recall	f1-score	support
0	0.67	0.74	0.70	329
1	0.65	0.57	0.61	281
accuracy			0.66	610
macro avg	0.66	0.66	0.66	610
weighted avg	0.66	0.66	0.66	610

```
confusion_matrix(y_train, pred_class_train)
```

executed in 9ms, finished 16:33:56 2021-11-28

```
array([[243,  86],
       [120, 161]], dtype=int64)
```

Classification report on testing data :

```
print(classification_report(y_test, pred_class_test))
```

executed in 19ms, finished 16:33:57 2021-11-28

	precision	recall	f1-score	support
0	0.65	0.76	0.70	142
1	0.65	0.52	0.57	120
accuracy			0.65	262
macro avg	0.65	0.64	0.64	262
weighted avg	0.65	0.65	0.64	262

```
confusion_matrix(y_test, pred_class_test)
```

executed in 9ms, finished 16:33:58 2021-11-28

```
array([[108,  34],
       [ 58,  62]], dtype=int64)
```

AUC and ROC for the training and testing data :

AUC for the Training Data: 0.731

AUC for the Test Data: 0.714




Figure-18

Comparison :

	LR Train	LR Test	LDA Train	LDA Test
<b>Accuracy</b>	0.66	0.66	0.66	0.65
<b>AUC</b>	0.73	0.72	0.73	0.71
<b>Recall</b>	0.57	0.52	0.57	0.52
<b>Precision</b>	0.66	0.66	0.65	0.65
<b>F1 Score</b>	0.61	0.58	0.61	0.57

Table-08

## 2.4 Inference: Based on these predictions, what are the insights and recommendations.

Please explain and summarise the various steps performed in this project. There should be proper business interpretation and actionable insights present.

- We have done predictions for both logistic regression and LDA. Since both results are the same.
- EDA analysis clearly indicates certain people aged above 50 are not interested much for holiday.
- People ranging from 35 to 45 generally opted for holiday.
- The important factor deciding the predictions are salary, age and educ.
- To improve holiday packages above 50 we can provide them religious destination places.
- For people earning more than 150000 we can provide them vacation holiday packages.
- Who have more than a number of older children we can provide packages in holiday vacation places.
- need to run promotional marketing campaigns or evaluate if we need to tie up with an alternate agency. It will increase sales.

